

Part_I_exploration

September 7, 2022

1 Part I - (Dataset Exploration Title)

1.1 by (Peter Lemule)

1.2 Introduction

Introduce the dataset: The dataset, 201902-fordgobike-tripdata.csv, is downloaded from Ford GoBike and licensed by Ford GoBike. This dataset includes 519,700 trips with 15 features such as locations, time, and user attributes. There are start and end stations. I noticed that most trips happen on the same stations, so I subset the dataset by choosing top 8 trips start stations with the most trips

1.3 Preliminary Wrangling

```
In [83]: # import all packages and set plots to be embedded inline
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb
import calendar
import math
import time

%matplotlib inline
```

Load in your dataset and describe its properties through the questions below. Try and motivate your exploration goals through this section.

```
In [84]: # Import Ford GoBike csv file into jupyter notebook
df = pd.read_csv('201902-fordgobike-tripdata.csv')
df.head()
```

```
Out[84]:
```

	duration_sec		start_time		end_time	\
0	52185	2019-02-28	17:32:10.1450	2019-03-01	08:01:55.9750	
1	42521	2019-02-28	18:53:21.7890	2019-03-01	06:42:03.0560	
2	61854	2019-02-28	12:13:13.2180	2019-03-01	05:24:08.1460	
3	36490	2019-02-28	17:54:26.0100	2019-03-01	04:02:36.8420	
4	1585	2019-02-28	23:54:18.5490	2019-03-01	00:20:44.0740	

	start_station_id	start_station_name	\
0	21.0	Montgomery St BART Station (Market St at 2nd St)	
1	23.0	The Embarcadero at Steuart St	
2	86.0	Market St at Dolores St	
3	375.0	Grove St at Masonic Ave	
4	7.0	Frank H Ogawa Plaza	

	start_station_latitude	start_station_longitude	end_station_id	\
0	37.789625	-122.400811	13.0	
1	37.791464	-122.391034	81.0	
2	37.769305	-122.426826	3.0	
3	37.774836	-122.446546	70.0	
4	37.804562	-122.271738	222.0	

	end_station_name	end_station_latitude	\
0	Commercial St at Montgomery St	37.794231	
1	Berry St at 4th St	37.775880	
2	Powell St BART Station (Market St at 4th St)	37.786375	
3	Central Ave at Fell St	37.773311	
4	10th Ave at E 15th St	37.792714	

	end_station_longitude	bike_id	user_type	member_birth_year	\
0	-122.402923	4902	Customer	1984.0	
1	-122.393170	2535	Customer	NaN	
2	-122.404904	5905	Customer	1972.0	
3	-122.444293	6638	Subscriber	1989.0	
4	-122.248780	4898	Subscriber	1974.0	

	member_gender	bike_share_for_all_trip
0	Male	No
1	NaN	No
2	Male	No
3	Other	No
4	Male	Yes

In [85]: df.shape

Out[85]: (183412, 16)

In [86]: df.describe()

Out[86]:

	duration_sec	start_station_id	start_station_latitude	\
count	183412.000000	183215.000000	183412.000000	
mean	726.078435	138.590427	37.771223	
std	1794.389780	111.778864	0.099581	
min	61.000000	3.000000	37.317298	
25%	325.000000	47.000000	37.770083	
50%	514.000000	104.000000	37.780760	

75%	796.000000	239.000000	37.797280
max	85444.000000	398.000000	37.880222

	start_station_longitude	end_station_id	end_station_latitude \
count	183412.000000	183215.000000	183412.000000
mean	-122.352664	136.249123	37.771427
std	0.117097	111.515131	0.099490
min	-122.453704	3.000000	37.317298
25%	-122.412408	44.000000	37.770407
50%	-122.398285	100.000000	37.781010
75%	-122.286533	235.000000	37.797320
max	-121.874119	398.000000	37.880222

	end_station_longitude	bike_id	member_birth_year
count	183412.000000	183412.000000	175147.000000
mean	-122.352250	4472.906375	1984.806437
std	0.116673	1664.383394	10.116689
min	-122.453704	11.000000	1878.000000
25%	-122.411726	3777.000000	1980.000000
50%	-122.398279	4958.000000	1987.000000
75%	-122.288045	5502.000000	1992.000000
max	-121.874119	6645.000000	2001.000000

In [87]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 183412 entries, 0 to 183411
Data columns (total 16 columns):
duration_sec      183412 non-null int64
start_time        183412 non-null object
end_time          183412 non-null object
start_station_id  183215 non-null float64
start_station_name 183215 non-null object
start_station_latitude 183412 non-null float64
start_station_longitude 183412 non-null float64
end_station_id    183215 non-null float64
end_station_name  183215 non-null object
end_station_latitude 183412 non-null float64
end_station_longitude 183412 non-null float64
bike_id           183412 non-null int64
user_type         183412 non-null object
member_birth_year 175147 non-null float64
member_gender     175147 non-null object
bike_share_for_all_trip 183412 non-null object
dtypes: float64(7), int64(2), object(7)
memory usage: 22.4+ MB
```

1.3.1 What is the structure of your dataset?

The datasets consists of 183412 trips and 16 features

1.3.2 What is/are the main feature(s) of interest in your dataset?

The Trip duration might be the main feature of interest of this dataset, as it's for sure affect the revenue of the company, that's why I will study on this analysis the effect of other factors like user_type, trip start and end times, gender and age on the trip duration.

1.3.3 What features in the dataset do you think will help support your investigation into your feature(s) of interest?

- user_type
- trip_duration
- end_station_name
- start_station_name
- start_time
- Gender and Age

```
In [88]: # checking for the null values
df.isnull().sum()
```

```
Out[88]: duration_sec          0
start_time                    0
end_time                      0
start_station_id             197
start_station_name           197
start_station_latitude        0
start_station_longitude       0
end_station_id               197
end_station_name             197
end_station_latitude          0
end_station_longitude         0
bike_id                      0
user_type                    0
member_birth_year            8265
member_gender                 8265
bike_share_for_all_trip       0
dtype: int64
```

```
In [89]: # Percentage of missing data
((df.isnull() | df.isna()).sum() * 100 / df.index.size).round(2)
```

```
Out[89]: duration_sec          0.00
start_time                    0.00
end_time                      0.00
start_station_id              0.11
```

```

start_station_name      0.11
start_station_latitude   0.00
start_station_longitude  0.00
end_station_id          0.11
end_station_name        0.11
end_station_latitude     0.00
end_station_longitude    0.00
bike_id                 0.00
user_type               0.00
member_birth_year       4.51
member_gender           4.51
bike_share_for_all_trip  0.00
dtype: float64

```

```

In [90]: # Since the we have that 4.51% is the highest record of missing values, it is better to
df.dropna(axis = 0, inplace = True)

```

```

In [91]: # Check the null values
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 174952 entries, 0 to 183411
Data columns (total 16 columns):
duration_sec          174952 non-null int64
start_time            174952 non-null object
end_time              174952 non-null object
start_station_id      174952 non-null float64
start_station_name    174952 non-null object
start_station_latitude 174952 non-null float64
start_station_longitude 174952 non-null float64
end_station_id        174952 non-null float64
end_station_name      174952 non-null object
end_station_latitude  174952 non-null float64
end_station_longitude 174952 non-null float64
bike_id               174952 non-null int64
user_type             174952 non-null object
member_birth_year     174952 non-null float64
member_gender         174952 non-null object
bike_share_for_all_trip 174952 non-null object
dtypes: float64(7), int64(2), object(7)
memory usage: 22.7+ MB

```

```

In [92]: # Check unique values
df.nunique()

```

```

Out[92]: duration_sec      4429
start_time      174941
end_time        174939

```

```

start_station_id      329
start_station_name     329
start_station_latitude 329
start_station_longitude 329
end_station_id         329
end_station_name       329
end_station_latitude   329
end_station_longitude   329
bike_id               4607
user_type              2
member_birth_year      75
member_gender          3
bike_share_for_all_trip 2
dtype: int64

```

```

In [93]: # Check for duplicated values
df.duplicated().sum()

```

```

Out[93]: 0

```

```

In [94]: # Check the datatypes
df.dtypes

```

```

Out[94]: duration_sec      int64
start_time                object
end_time                  object
start_station_id          float64
start_station_name        object
start_station_latitude     float64
start_station_longitude    float64
end_station_id            float64
end_station_name          object
end_station_latitude       float64
end_station_longitude      float64
bike_id                   int64
user_type                  object
member_birth_year          float64
member_gender              object
bike_share_for_all_trip    object
dtype: object

```

```

In [95]: # columns start_time and end_time should be in the datetime format
# columns user_type and bike_share_for_all_trip should be in categorical datatype and n
# start_station_id, end_station_id and bike_id should be object data type not numerical
df.start_time = pd.to_datetime(df.start_time)
df.end_time = pd.to_datetime(df.end_time)
df['user_type'] = df['user_type'].astype('category')
df['bike_share_for_all_trip'] = df['bike_share_for_all_trip'].astype('category')
df['start_station_id'] = df['start_station_id'].astype('object')

```

```
df['end_station_id'] = df['end_station_id'].astype('object')
df['bike_id'] = df['bike_id'].astype('object')
```

```
In [96]: # Corfirm the changes made on the dtype
df.dtypes
```

```
Out[96]: duration_sec          int64
start_time          datetime64[ns]
end_time            datetime64[ns]
start_station_id      object
start_station_name    object
start_station_latitude float64
start_station_longitude float64
end_station_id        object
end_station_name      object
end_station_latitude  float64
end_station_longitude float64
bike_id              object
user_type             category
member_birth_year     float64
member_gender         object
bike_share_for_all_trip category
dtype: object
```

```
In [99]: # Extract start_time_month and dayofweek, information from the start_time
df['start_time_dayofweek'] = df['start_time'].dt.strftime('%a')
df['start_time_month'] = df['start_time'].dt.strftime('%B')
```

```
In [100]: # day categories in order
weekday = ['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun']
# change type to ordered categorical
df['start_time_dayofweek'] = pd.Categorical(df['start_time_dayofweek'], categories=weekday, ordered=True)

# check counts and order (not required)
df['start_time_dayofweek'].value_counts().sort_index()
```

```
Out[100]: Mon      25641
Tue       30584
Wed       28426
Thu       33712
Fri       27663
Sat       14414
Sun       14512
Name: start_time_dayofweek, dtype: int64
```

```
In [101]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 174952 entries, 0 to 183411
```

Data columns (total 18 columns):

```
duration_sec          174952 non-null int64
start_time            174952 non-null datetime64[ns]
end_time              174952 non-null datetime64[ns]
start_station_id      174952 non-null object
start_station_name    174952 non-null object
start_station_latitude 174952 non-null float64
start_station_longitude 174952 non-null float64
end_station_id        174952 non-null object
end_station_name      174952 non-null object
end_station_latitude  174952 non-null float64
end_station_longitude 174952 non-null float64
bike_id               174952 non-null object
user_type             174952 non-null category
member_birth_year     174952 non-null float64
member_gender         174952 non-null object
bike_share_for_all_trip 174952 non-null category
start_time_dayofweek  174952 non-null category
start_time_month      174952 non-null object
dtypes: category(3), datetime64[ns](2), float64(5), int64(1), object(7)
memory usage: 21.9+ MB
```

```
In [102]: # let's change the duration_sec to duration_min to make it more compact
df['duration_min'] = df['duration_sec'] / 60
```

```
In [103]: df.duration_min.describe()
```

```
Out[103]: count      174952.000000
          mean         11.733379
          std         27.370082
          min          1.016667
          25%          5.383333
          50%          8.500000
          75%         13.150000
          max         1409.133333
          Name: duration_min, dtype: float64
```

```
In [104]: # Looks like for at least 75% of the data are less than one hour, so end_time_hour will
          # Extract Start_time_hour from start_time information
df['start_time_hour'] = df['start_time'].dt.hour
```

```
In [105]: # checking the unique values of the start_time_month
df['start_time_month'].unique()
```

```
Out[105]: array(['February'], dtype=object)
```

```
In [106]: # create a new column member_age which is easier to relate with that the birth year th
          # since the data set is 2019, we will be subtracting the birth year from 2019
df['member_age'] = 2019 - df['member_birth_year']
df['member_age'].head(2)
```



```
Out[106]: 0    35.0
          2    47.0
          Name: member_age, dtype: float64
```

```
In [107]: # Changing the dtype of both member_birth_year and member_age from float to int datatype
          df['member_age'] = df['member_age'].astype(int)
          df['member_birth_year'] = df['member_birth_year'].astype(int)
```

```
In [108]: print(df['member_age'].head(2))
          print(df['member_birth_year'].head(2))
```

```
0    35
2    47
Name: member_age, dtype: int64
0    1984
2    1972
Name: member_birth_year, dtype: int64
```

```
In [109]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 174952 entries, 0 to 183411
Data columns (total 21 columns):
duration_sec          174952 non-null int64
start_time            174952 non-null datetime64[ns]
end_time              174952 non-null datetime64[ns]
start_station_id      174952 non-null object
start_station_name     174952 non-null object
start_station_latitude 174952 non-null float64
start_station_longitude 174952 non-null float64
end_station_id        174952 non-null object
end_station_name       174952 non-null object
end_station_latitude   174952 non-null float64
end_station_longitude  174952 non-null float64
bike_id               174952 non-null object
user_type             174952 non-null category
member_birth_year      174952 non-null int64
member_gender          174952 non-null object
bike_share_for_all_trip 174952 non-null category
start_time_dayofweek   174952 non-null category
start_time_month       174952 non-null object
duration_min          174952 non-null float64
start_time_hour        174952 non-null int64
member_age             174952 non-null int64
dtypes: category(3), datetime64[ns](2), float64(5), int64(4), object(7)
memory usage: 25.9+ MB
```

```
In [110]: quantile_range = list(np.arange(0.1, 1, 0.05))
          df['duration_min'].quantile(quantile_range)
```

```
Out[110]: 0.10      3.550000
          0.15      4.200000
          0.20      4.800000
          0.25      5.383333
          0.30      5.950000
          0.35      6.550000
          0.40      7.183333
          0.45      7.816667
          0.50      8.500000
          0.55      9.250000
          0.60     10.050000
          0.65     10.950000
          0.70     11.966667
          0.75     13.150000
          0.80     14.616667
          0.85     16.500000
          0.90     19.350000
          0.95     25.516667
          Name: duration_min, dtype: float64
```

```
In [111]: # We remove outliers row
          df = df[df['duration_min'] <= df['duration_min'].quantile(0.99)]
```

```
In [112]: df.describe()
```

```
Out[112]:
```

	duration_sec	start_station_latitude	start_station_longitude	\
count	173204.000000	173204.000000	173204.000000	
mean	612.726138	37.771176	-122.351553	
std	425.821242	0.100544	0.117839	
min	61.000000	37.317298	-122.453704	
25%	321.000000	37.770407	-122.411901	
50%	506.000000	37.780760	-122.398279	
75%	777.000000	37.797320	-122.283093	
max	3176.000000	37.880222	-121.874119	

	end_station_latitude	end_station_longitude	member_birth_year	\
count	173204.000000	173204.000000	173204.000000	
mean	37.771362	-122.351117	1984.808665	
std	0.100439	0.117387	10.112763	
min	37.317298	-122.453704	1878.000000	
25%	37.770407	-122.411647	1980.000000	
50%	37.781010	-122.397405	1987.000000	
75%	37.797673	-122.285171	1992.000000	
max	37.880222	-121.874119	2001.000000	

	duration_min	start_time_hour	member_age
--	--------------	-----------------	------------

count	173204.000000	173204.000000	173204.000000
mean	10.212102	13.454822	34.191335
std	7.097021	4.739169	10.112763
min	1.016667	0.000000	18.000000
25%	5.350000	9.000000	27.000000
50%	8.433333	14.000000	32.000000
75%	12.950000	17.000000	39.000000
max	52.933333	23.000000	141.000000

```
In [113]: # copy the data set
df_clean = df.copy()
df_clean.to_csv('clean_ford_data.csv')
```

```
In [114]: df.describe()
```

```
Out[114]:
```

	duration_sec	start_station_latitude	start_station_longitude	\
count	173204.000000	173204.000000	173204.000000	
mean	612.726138	37.771176	-122.351553	
std	425.821242	0.100544	0.117839	
min	61.000000	37.317298	-122.453704	
25%	321.000000	37.770407	-122.411901	
50%	506.000000	37.780760	-122.398279	
75%	777.000000	37.797320	-122.283093	
max	3176.000000	37.880222	-121.874119	

	end_station_latitude	end_station_longitude	member_birth_year	\
count	173204.000000	173204.000000	173204.000000	
mean	37.771362	-122.351117	1984.808665	
std	0.100439	0.117387	10.112763	
min	37.317298	-122.453704	1878.000000	
25%	37.770407	-122.411647	1980.000000	
50%	37.781010	-122.397405	1987.000000	
75%	37.797673	-122.285171	1992.000000	
max	37.880222	-121.874119	2001.000000	

	duration_min	start_time_hour	member_age
count	173204.000000	173204.000000	173204.000000
mean	10.212102	13.454822	34.191335
std	7.097021	4.739169	10.112763
min	1.016667	0.000000	18.000000
25%	5.350000	9.000000	27.000000
50%	8.433333	14.000000	32.000000
75%	12.950000	17.000000	39.000000
max	52.933333	23.000000	141.000000

1.4 Univariate Exploration


duration_min

```

In [115]: # plotting the duration_min data on a normal scale
binsize = 5
bins = np.arange(0, df['duration_min'].max()+binsize, binsize)

plt.figure(figsize=[10, 5])
plt.hist(data = df, x = 'duration_min', bins = bins)
plt.title('Distribution of Trip Durations')
plt.xlabel('Duration in Minutes')
plt.ylabel('Count of Trips')
plt.axis([-10, 200, 0, 70000])
plt.show()

```




```

In [116]: log_binsize = 0.025
bins = 10 ** np.arange(0, np.log10(df['duration_min'].max())+log_binsize, log_binsize)

plt.figure(figsize=[10, 7]);
plt.hist(data = df, x = 'duration_min', bins = bins);
plt.xscale('log');
plt.xticks([1, 2, 5, 10, 20, 50, 100, 200, 500], [1, 2, 5, 10, 20, 50, 100, 200, 500]);
plt.xlabel('Duration_min');


```



```
In [117]: # Functions to avoid repetition of codes
def countplot(data, x, color, order = None):
    plt.figure(figsize=[10, 7])
    base_color = sb.color_palette()[color]
    if order:
        order = df[x].value_counts().index
    sb.countplot(data = data, x = x, color = base_color, order = order)
    plt.title('Distribution of ' + x)
    plt.xlabel(x)
    plt.ylabel('Counts');

def piechart(x):
    sorted_counts = df[x].value_counts()
    plt.figure(figsize=[10,7])
    plt.pie(sorted_counts, labels = sorted_counts.index, startangle = 90, counterclockwise=True)
    plt.axis('square')
    plt.title('Pie Chart of ' + x);

In [118]: # a bar plot showing the peak start hours
countplot(data = df, x = 'start_time_hour', color = 5)
```




Check for start hours with the peak time

member_gender


```
In [119]: df['member_gender'].value_counts()
```

```
Out[119]: Male      129347  
         Female    40310  
         Other      3547  
         Name: member_gender, dtype: int64
```

```
In [120]: # Bar plot showing that most Bikers are Male  
          countplot(data = df, x = 'member_gender', color=9)
```



```
In [121]: # A pie chart to justify the bar chart that most Bikers are Male  
          piechart('member_gender')
```



start and end stations

```
In [122]: df['start_station_name'].value_counts()[1:10]
```

```
Out[122]: San Francisco Caltrain Station 2 (Townsend St at 4th St)    3399
Berry St at 4th St                                                    2933
Montgomery St BART Station (Market St at 2nd St)                    2694
Powell St BART Station (Market St at 4th St)                          2570
San Francisco Caltrain (Townsend St at 4th St)                       2565
San Francisco Ferry Building (Harry Bridges Plaza)                   2493
Howard St at Beale St                                                 2211
Steuart St at Market St                                               2174
Powell St BART Station (Market St at 5th St)                          2098
Name: start_station_name, dtype: int64
```



```
In [123]: df['end_station_name'].value_counts()[1:10]
```

```
Out[123]: Market St at 10th St                3697
Montgomery St BART Station (Market St at 2nd St)  3438
San Francisco Ferry Building (Harry Bridges Plaza)  3109
San Francisco Caltrain (Townsend St at 4th St)    2864
Powell St BART Station (Market St at 4th St)      2824
Berry St at 4th St                               2772
The Embarcadero at Sansome St                    2283
Steuart St at Market St                          2256
Powell St BART Station (Market St at 5th St)      2122
Name: end_station_name, dtype: int64
```

```
In [124]: # since most start and end stations have the same frequency, we then add both together
# select top 10
start_stations = df['start_station_name'].value_counts()
end_stations = df['end_station_name'].value_counts()
station_total = start_stations + end_stations
```

```
In [125]: station_total.head()
```

```
Out[125]: 10th Ave at E 15th St                90
10th St at Fallon St                        683
10th St at University Ave                  452
11th St at Bryant St                      1733
11th St at Natoma St                      1634
dtype: int64
```

```
In [126]: station_total = pd.DataFrame(station_total, index=None).reset_index().rename(columns=

station_total.head()
```

```
Out[126]:
```

	station	count
0	10th Ave at E 15th St	90
1	10th St at Fallon St	683
2	10th St at University Ave	452
3	11th St at Bryant St	1733
4	11th St at Natoma St	1634

```
In [127]: top_10 = pd.DataFrame(start_stations.sort_values(ascending=False)[:10],
                                index=None).reset_index().rename(columns={'index': 'station', 'sta

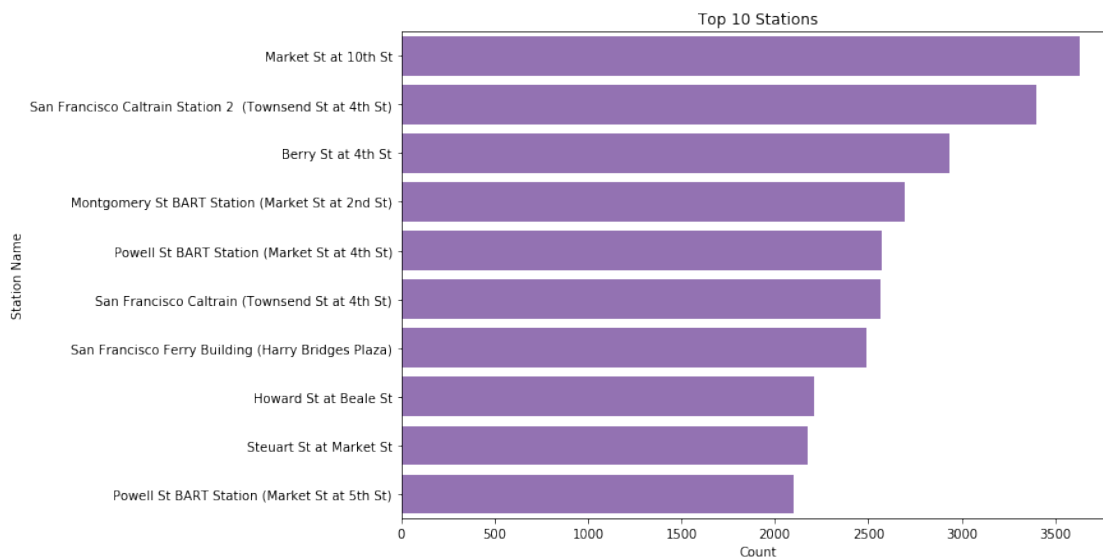
In [128]: top_10
```

```
Out[128]:
```

	station	count
0	Market St at 10th St	3634
1	San Francisco Caltrain Station 2 (Townsend St...	3399
2	Berry St at 4th St	2933
3	Montgomery St BART Station (Market St at 2nd St)	2694

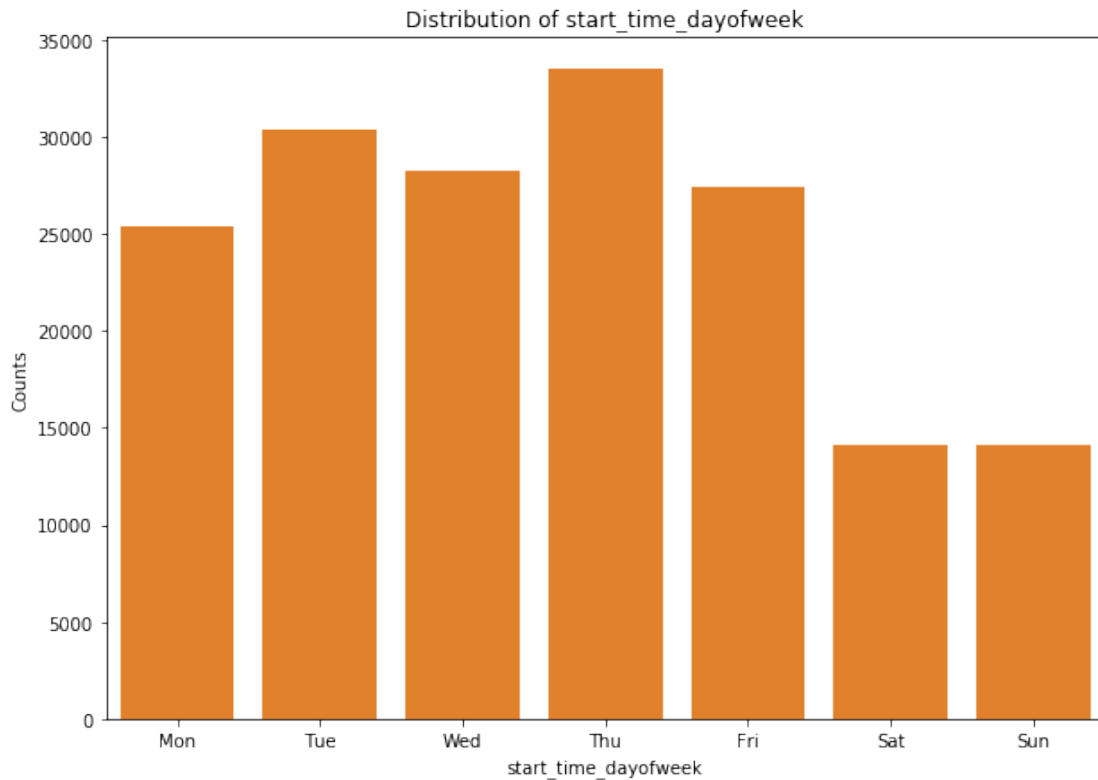
4	Powell St BART Station (Market St at 4th St)	2570
5	San Francisco Caltrain (Townsend St at 4th St)	2565
6	San Francisco Ferry Building (Harry Bridges Pl...	2493
7	Howard St at Beale St	2211
8	Steuart St at Market St	2174
9	Powell St BART Station (Market St at 5th St)	2098

```
In [129]: # plot the top ten stations
plt.figure(figsize = [10,7])
base_color = sb.color_palette()[4]
sb.barplot(data=top_10, y = 'station', x='count', color=base_color)
plt.title('Top 10 Stations')
plt.xlabel('Count')
plt.ylabel('Station Name')
plt.show();
```



Days of the week

```
In [130]: # A bar plot to show the counts for days of the week
# Most bikers work during the weekdays
countplot(data = df, x = 'start_time_dayofweek', color = 1)
```




Age


```
In [131]: df['member_age'].describe()
```

```
Out[131]: count    173204.000000
          mean       34.191335
          std        10.112763
          min        18.000000
          25%        27.000000
          50%        32.000000
          75%        39.000000
          max        141.000000
          Name: member_age, dtype: float64
```

```
In [132]: sb.boxplot(data = df, y = 'member_age');
```



```
In [133]: # This plot shows the age range of the Bikers
binsize = 2
bins = np.arange(16, df['member_age'].max()+binsize, binsize)
plt.hist(data = df, x = 'member_age', bins = bins);
plt.title("Distribution of Users' Age")
plt.xlabel('Age')
plt.ylabel('Number of Users')
plt.show()
```



```
In [134]: # group the ages into generation
def age_distribution(yob):
    if yob >= 1928 and yob <= 1945:
        return('Post War')
    elif yob >= 1946 and yob <= 1964:
        return('Baby Boomers')
    elif yob >= 1965 and yob <= 1980:
        return('Gen X')
    elif yob >= 1981 and yob <= 1996:
        return('Millenials')
    elif yob >= 1997 and yob <= 2012:
        return('Gen Z')
    else: np.NaN

In [135]: # create member generation variables
df['member_generation'] = df['member_birth_year'].apply(age_distribution)

# create generations variable
generations = ['Post War', 'Baby Boomers', 'Gen X', 'Millenials', 'Gen Z']

# Order member_generation categorically
df['member_generation'] = pd.Categorical(df['member_generation'], categories = generat
```

```
# Counts in order
df['member_generation'].value_counts().sort_index()
```

```
/opt/conda/lib/python3.6/site-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

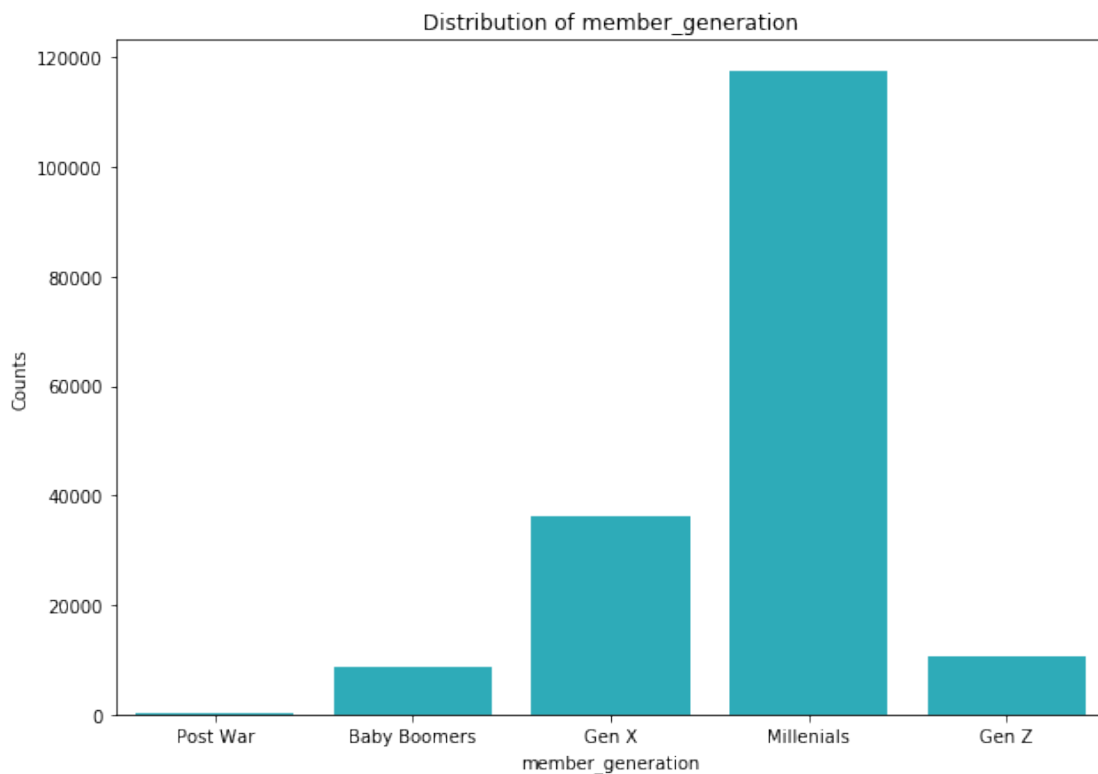
See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>

```
/opt/conda/lib/python3.6/site-packages/ipykernel_launcher.py:8: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>


```
Out[135]: Post War      289
          Baby Boomers  8592
          Gen X        36189
          Millenials   117290
          Gen Z        10768
          Name: member_generation, dtype: int64
```

```
In [136]: countplot(data = df, x = 'member_generation', color = 9)
```



Bikers Type

```
In [137]: # Plot a pie chart to show the user_type  
          # It shows that most of the Bikers are Subscriber  
          piechart('user_type')
```



Plotting the duration of trips using histogram shows that we need to scale. I apply log scaling to show to the plot which now shows that the duration of most trips are between 6 - 15 mins

The rides across the week clearly shows that Thursday has the highest number of trips. Weekdays seems to be the most significant counts, while weekends (Saturday and Sunday) have the least counts

The start stations and end stations have similar top records which is why it was added together and it was plotted in the graph to show the most top 10 most frequent places in San Francisco

The Bikers were splitted generations using their birth year after plotting it was observed that the millenials generation has the highest number i.e Bikers that are born between 1981 and 1996. The generation was gotten using birth_year

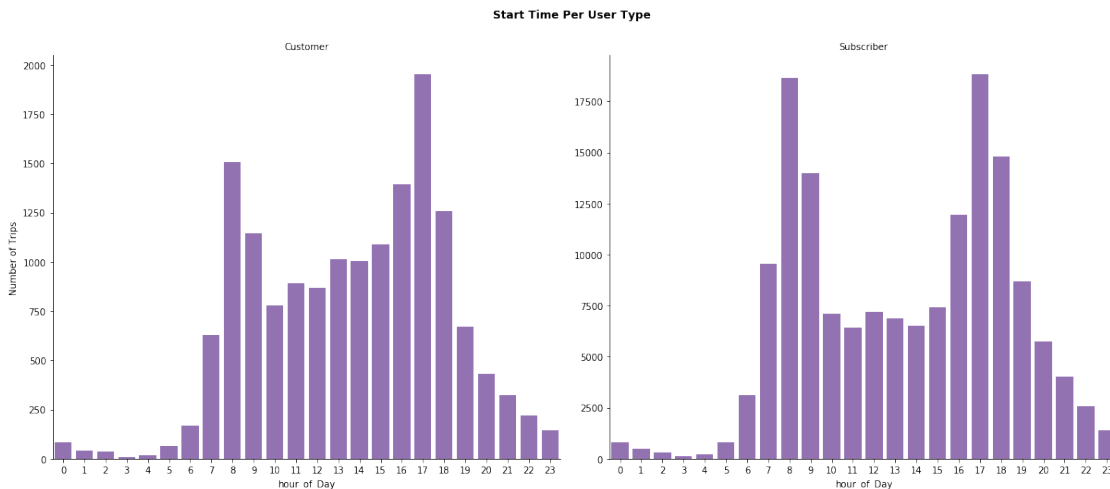
Majority of the Bikers are subscribers, which can be seen from the pie chart

The hour of the day and the day of the week were extracted from the timestamp.

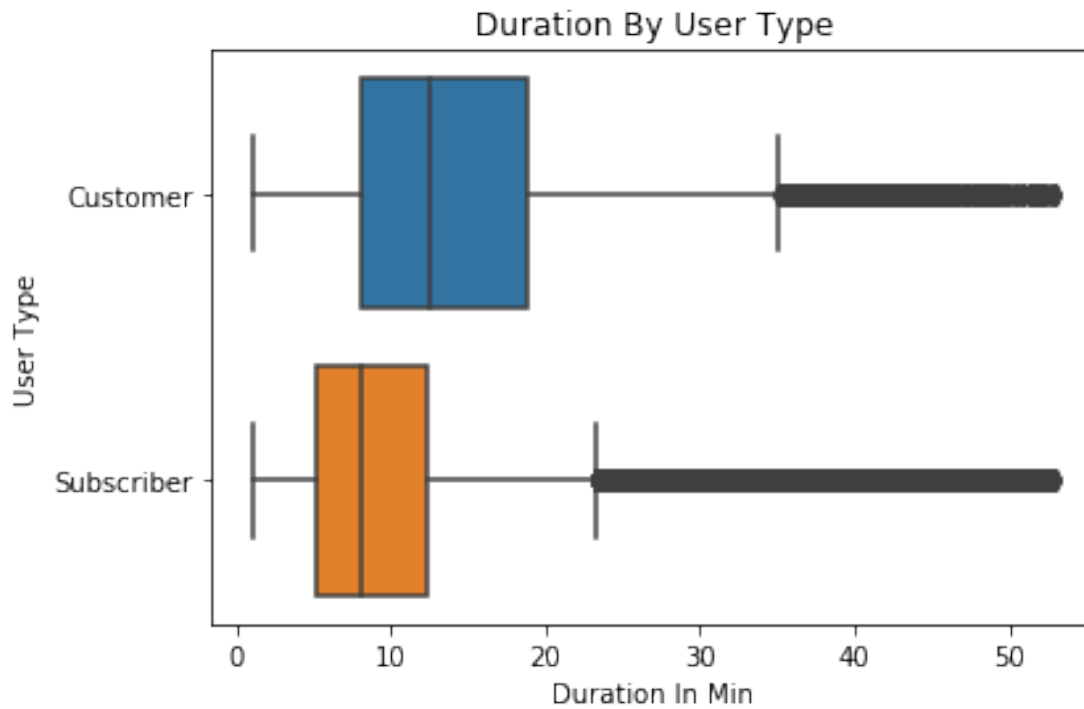
1.5 Bivariate Exploration

In this section, investigate relationships between pairs of variables in your data. Make sure the variables that you cover here have been introduced in some fashion in the previous section (univariate exploration).


```
In [138]: # User type per hour
base_color = sb.color_palette()[4]
user_per_hour = sb.factorplot(data=df, x='start_time_hour', col="user_type", kind='count',
                              color = base_color, size=7, aspect=1.2)
user_per_hour.set_axis_labels("hour_of_Day", "Number of Trips")
user_per_hour.set_titles("{col_name}")
user_per_hour.fig.suptitle('Start Time Per User Type', y=1.05, fontsize=12, fontweight='bold')
```



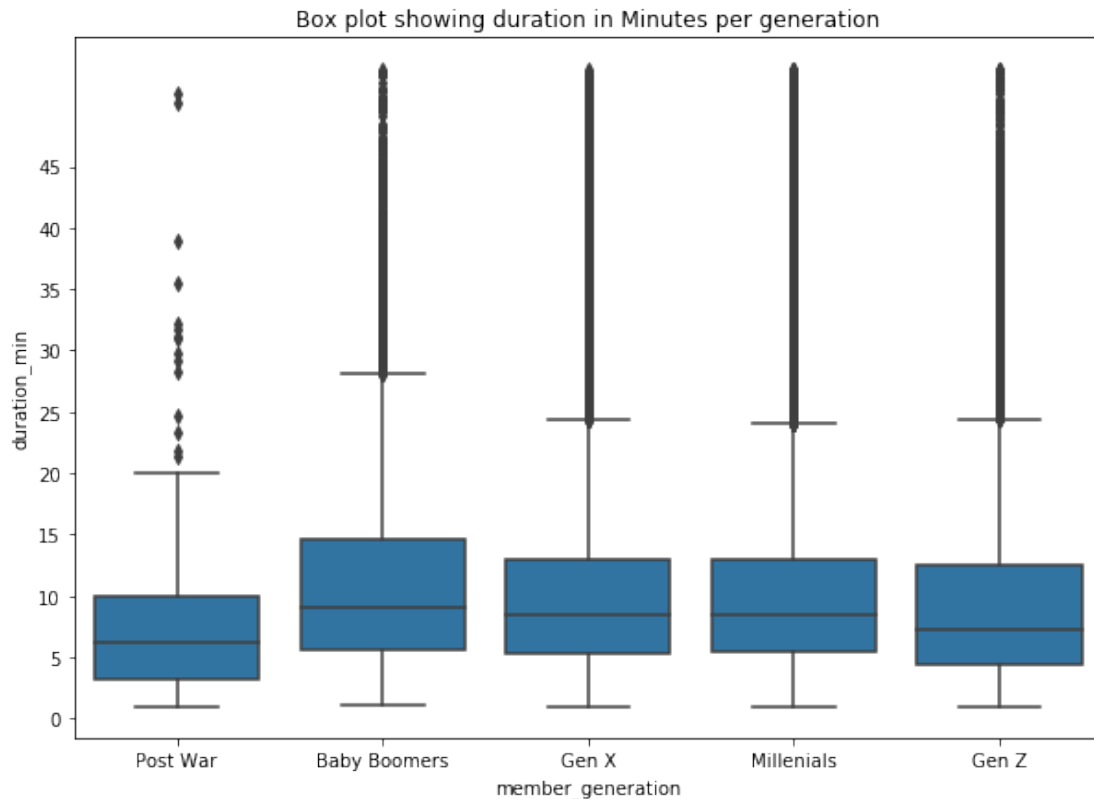

```
In [139]: # Duration of usertypes per minutes
sb.boxplot(data =df , y='user_type', x='duration_min')
plt.title('Duration By User Type')
plt.xlabel('Duration In Min')
plt.ylabel('User Type');
```



```
In [140]: base_color = sb.color_palette()[2]
plt.figure(figsize=(10,7))
sb.violinplot(data = df, x='member_generation', y= 'duration_min', color = base_color)
plt.title('Violin plot showing duration in Minutes per generation')
plt.yticks(np.arange(0, 50, step=5));
```



```
In [141]: base_color = sb.color_palette()[0]
plt.figure(figsize=(10,7))
sb.boxplot(data = df, x='member_generation', y= 'duration_min', color = base_color)
plt.title('Box plot showing duration in Minutes per generation')
plt.yticks(np.arange(0, 50, step=5));
```



From the factor plot, after splitting the hour per usage into customer and subscribers, it clearly shows that most customers have their start time at 17:00, followed by 8:00 but subscribers have thier peake hours at 8:00 and 17:00 which clearly shows that some of the importance of Bivariate to Univariate

I found that the duration minutes were different for Customers and Subscribers. Subscribers have their most duration minutes between 7 - 16 minutes while Customers have their most duration between 5 - 12 minutes.

The box plot gives more summary details visibly than the violin plot. From the box plot, it is noticed that the Baby Boomers ride for the longest time compared to the other generations. The Post War generation have the least riding duration and this is likely due to age.

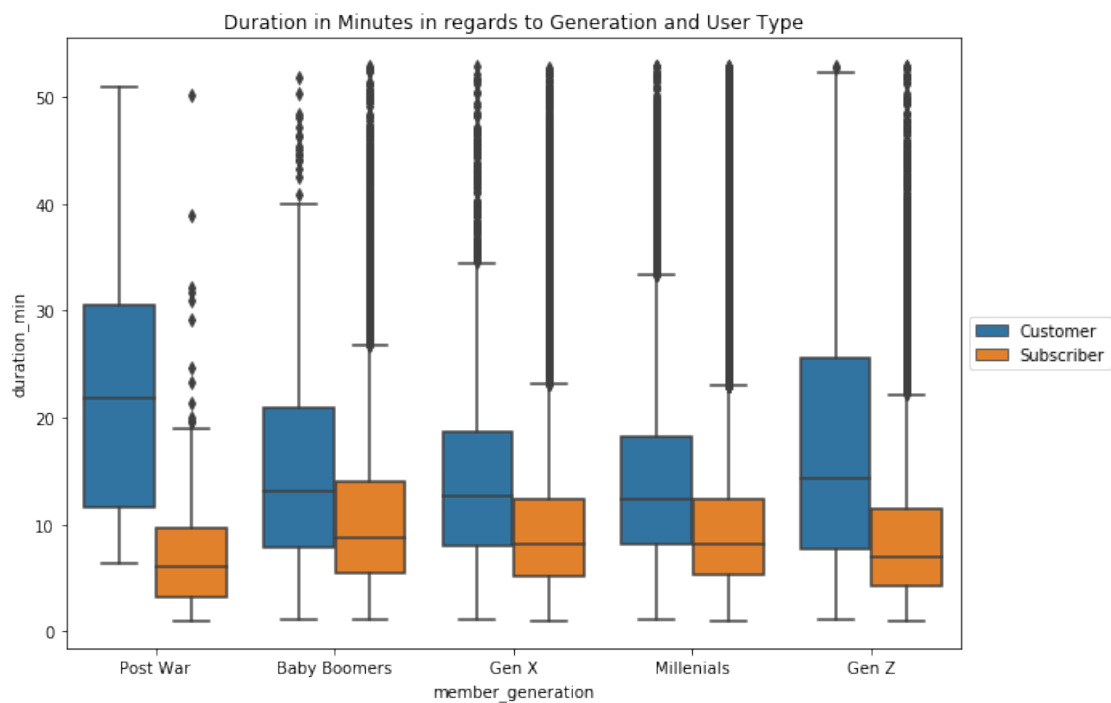
1.6 Multivariate Exploration

```
In [142]: # This plot compares the start_time_hour and the duration_min in regards to generation
g = sb.FacetGrid(data = df, col = 'member_generation', col_wrap = 3)
g.map(sb.stripplot, 'start_time_hour', 'duration_min', size = 2, jitter = 0.35, order

g.fig.set_size_inches(24, 14)
```



```
In [143]: plt.figure(figsize=[10, 7])
          sb.boxplot(data = df, x = 'member_generation', y = 'duration_min', hue = 'user_type')
          plt.legend(loc = 6, bbox_to_anchor = (1.0, 0.5))
          plt.title('Duration in Minutes in regards to Generation and User Type');
```



From the boxplot for duration in Minutes in regards to Generation and User Type shows that it's the post customer usertype that has the least duration minutes and the post war subscribers have the Highest duration in Minutes

1.7 Conclusions

From the data set, 8am and 9pm have the highest number of trips

Customers have the lowest duration minutes while subscribers are the highest

Majority of the Bikers are subscribers, which can be seen from the pie chart

Most bikers falls within the Millenials Generation that is they are born between 1981 and 1996